

Cordwell Eval Games II: Release Day

Audit the Numbers. Design the Gauntlet. Survive the Attack.

How the Game Works

It is release day at Cordwell Home and Hardware. A new version of the store-associate assistant is waiting to ship, the weekend sale starts Saturday, and your engineering manager is already drafting the announcement. Three rounds, thirty to forty-five minutes, points on the board.

Round 1: Ship It or Block It

You are the release review board. Below are the MLflow run cards for the current production version (v2) and the candidate (v3), exactly as logged. Then comes the message from your engineering manager. Your job: decide whether these two runs actually prove what the message claims.

Run card: cordwell_qa_v2 (July 1)

Params	
app_version	v2
system_prompt	"Answer the question using the context. If uncertain, say so."
retrieval_k	3
backend_mode	offline (deterministic judge callables)
model_tag	not applicable in offline mode
Tags	
seed_query_set	q_set_2026_07_01 (12 queries, locked at authoring time)
Metrics	
groundedness_mean	0.90
answer_relevance_mean	0.85
Artifacts	
files	cordwell_records_v2.csv

Run card: cordwell_qa_v3 (July 28)

Params	
app_version	v3
system_prompt	"Answer using only the context. Cite the aisle number when available. If uncertain, say so."
retrieval_k	5
backend_mode	ollama (live judge)

model_tag	gemma4
Tags	
seed_query_set	q_set_2026_07_28 (12 queries; 3 replaced to remove discontinued products)
Metrics	
groundedness_mean	0.92
answer_relevance_mean	0.84
Artifacts	
files	cordwell_records_v3.csv, triad_correlation.png

Message from your engineering manager: “Groundedness is up two points. Ship v3 before the weekend sale.”

Task 1: The audit (1 point per problem found)

Every number on those cards is real. The comparison may still be broken. Find every reason these two runs cannot be fairly compared. There are at least FOUR planted problems. Read every line; the params are evidence.

Problem 1: _____

Problem 2: _____

Problem 3: _____

Problem 4: _____

Anything else: _____

Task 2: The verdict

Circle one as a board, and be ready to defend it to the room:

SHIP IT BLOCK IT RERUN IT

Task 3: The Run Protocol (1 point)

Write the rules that would make the rerun decisive. What must be held constant, what may change per run, and what must be logged so the next board can audit the comparison without asking anyone?

Round 2: Design the Gauntlet

v3 goes back to the lab under your protocol. Meanwhile, Cordwell is standing up new AI systems, and each needs an evaluation strategy before it writes a single production sentence. Each team picks ONE of these systems (and each team must cover something different):

System	What it does
Product Q&A RAG	The store-associate assistant answering product and aisle questions for consumers
Marketing Copy Generator	Writes promotional copy for the weekly flyer and product pages
Internal Code Assistant	Answers engineers' questions grounded in Cordwell's internal codebase docs
Review Summarizer	Condenses customer reviews into a short summary on each product page

The Strategy Canvas (2 points)

Fill in every row for YOUR system. Generic answers score nothing; the canvas must fit the system you drew. Thresholds are policy choices, so name an owner and a rationale, not just a number. Prepare a 60-second pitch.

- **Triad legs:** which of the three apply to your system, and why? Dropping a leg requires a reason.

- **Two custom checks:** risks the Triad cannot see. Anything a rubric can express can become a metric.

- **Thresholds:** one example gate, who owns the number, and why that strictness.

- **CI gate (every commit):** exactly what runs, and what makes it fast and deterministic.

- **Nightly:** what runs here instead of CI, and why it can afford to.

- **Production monitoring:** one signal you would watch, and the condition that pages someone.

Round 3: Break the Bot

Your gauntlet is designed. Time to find out if it stops anything. Write TWO red-team probes against the system of the team clockwise from you, drawn from two DIFFERENT probe families:

- **Prompt injection:** input that tries to override the system's instructions.
- **Jailbreak:** a persona or authority trick that tries to unlock something restricted.
- **Hallucination trigger:** a question about something that does not exist, inviting a confident invention.

- **Bias probe:** paired inputs that differ only in framing or demographics, checking whether quality or tone shifts.

Rule: each probe must be a plausible input a real user could send to that system. Jokes with no domain contact score nothing.

Our attack (1 point for exposing a real gap)

Target system: _____

Probe 1 (family:):

Probe 2 (family:):

Our defense (1 point for naming the mechanism)

For each probe you receive: name the specific pipeline stage or the specific check on your canvas that catches it or admit the gap and name the fix you would add. “Our eval would catch that” scores nothing; “our seeded adversarial set plus the no-invented-promotions check fails it in the offline suite” scores.

Incoming probe 1, caught by (or gap plus fix):

Incoming probe 2, caught by (or gap plus fix):
